
Collaborative Filtering Recommendation System for Steam with Deep and Variational Autoencoders

Daniel Liang^{* 1} Shuidie Yao^{* 1}

Abstract

Recommendation systems are commonly used to provide interesting and meaningful content personalized to a users preferences. While multiple approaches are used to create them, autoencoders are commonly used as the model for collaborative filtering recommendation systems. We prepare data from a Steam dataset into a matrix based on the time that users have played specific games. Additional matrices containing game names, game genres and user reviews from individual users are used as latent variables in our models. We build recommendation systems based on the formatted data, two with deep autoencoder models and two with variational autoencoder models. Finally, we examine the differences between the various aspects of the models.

1. Introduction

The goal of recommendation systems is to recommend items for customers based on their past reflection of purchased and/or used goods. To deal with a large amount of data, two primary types of algorithms are prevalent when creating recommendation systems: content-based filtering algorithms and collaborative filtering algorithms. Content-based filtering algorithms produce recommendations by looking at past positive reactions from a user to a certain item and recommending items based on similar attributes. This could take the form of recommending another popular action movie because the user had liked a similar action movie in the past. Collaborative filtering algorithms differ because they recommend based on multiple users historical preferences on a set of items to recommend items from similar users rather than rely on item attributes.

The data for collaborative filtering generally has a spe-

cific pattern to consider. A single customer usually buys a few types of items out of the hundreds and thousands of total items offered. These items tend to be popular items that other users own. Most items therefore are favored only by a few people, while top items are extremely popular in comparison. This popularity is generally shown in a form that indicates what the user thinks of each item. In the case of television shows, this would most likely be the amount of time a user spends watching each show. In a typical representation of such data—a two-dimensional matrix where rows are users, columns are items, and each element is a score of one user’s preference of one item—we would see that a majority of the total shows offered would have a watch time of zero, since most users don’t interact with most items.

Neural networks have been very effective in investigating large and sparse matrices. In particular, a popular method in building recommendation systems to describe features contained in these types of matrices is to use an autoencoder. Autoencoders consist of two neural networks, where one is an encoder that trains a function to bring high dimension input data to a bottleneck layer where the number of neurons is the smallest. The other neural network is a decoder that converts the encoded input back to its original dimension. The autoencoder works well by learning to extract the most relevant features in the bottleneck from the latent space, and the latent space contains the only information a decoder can use to reconstruct the input as faithfully as possible.

Variational autoencoders share roughly the same process but train different things. It assumes that the data has a certain distribution, and finds the parameters of the distribution during training. Despite the reconstruction loss, its loss function also contains a regularizer, which is the Kullback-Leibler divergence between the encoder distribution and the decoder distribution. Variational autoencoders implicitly create a probability distribution based on the data it is given.

Autoencoders have been successfully implemented for various different datasets and have produced impressive results. A deep autoencoder researched by NVIDIA on the Netflix prize dataset that outperforms previously built autoencoder models (Kuchaiev & Ginsburg, 2017). Similarly, a variational autoencoder was implemented on the Movie-

^{*}Equal contribution ¹Courant Institute of Mathematical Sciences, New York University, New York City, United States. Correspondence to: Daniel Liang <dzl217@nyu.edu>, Shuidie Yao <sy2327@nyu.edu>.

Lens 20M dataset (Liang et al., 2018). The implementations discussed in this paper are based on the models proposed by these authors for deep and variational autoencoders on a Steam dataset.

2. Method

Four recommendation systems are built using one procedure but with different autoencoders. The data we use required cleaning and formatting prior to training. We train the models on Google Colaboratory. Recommendations were produced as well as data from the models' training that allows us to compare the performance of the two autoencoders. The full code can be found at <http://github.com/yshuidie/steamNNRecom>.

2.1. Data and Preparation

We use the Steam video game dataset from the collection of recommendation system datasets offered by Julian McAuley from the University of California San Diego. The data used can be found and downloaded at <http://cseweb.ucsd.edu/~jmcauley/datasets.html>. On the Steam platform users can buy games, play those games and write reviews of them. Based on our experience as players along with our observations by profiling the data, we see that most players play less than half of the games they own. This shows that some players are generally interested in the games they purchase, but have preferences towards the games they have played. Additionally, players usually spend a significant amount of time on their favorite games more than the other games they own. We therefore decide to treat a user's play time for a certain game as a scoring metric to show how much a user likes that game. From the game play time of each user, we formulate a user preference matrix.

We denote the play time matrix T , where each entry T_{ij} represents the play time in minutes that user i spent on game j . We create an approximation of a probability matrix S , which represents the user preference matrix, that a player would play a game with the following properties:

- A function exists that maps values $T_{ij} \in (0, \infty)$ to $S_{ij} \in (0, 1)$
- The probability S_{ij} increases with playtime T_{ij}
- S_{ij} and T_{ij} have a non-linear relationship. S_{ij} should increase marginally as T_{ij} increases.

According to these needs, we use the tanh function to create the new probability matrix. We set the maximum playtime of any game to 50 hours, and based on the observation that $\tanh(2) \approx 0.96 \approx 1$, we construct the probability

matrix such that

$$S_{ij} = \tanh\left(\frac{T_{ij} * 2}{50 * 60}\right)$$

We perform several tasks in order to properly prepare the data for training. First, we filter out games that were never played by any user and users who never played any games from our Pandas dataframe. Next, due to memory limits within Google Colaboratory, we limit the number of games we'd investigate. We cap the game ID that steam uses for game identification and consider all games with an ID lower than 10000. At a result, 354 games are considered. For each user, we extract a list of games that the user played, the names of those games, the games genres and the reviews the user made. The words in this data are transformed into integers by `keras.one_hot()`, so that they could be fed into the autoencoders. In addition, we reformat the game information in a new dataframe to include all the games with information about their game price, developer and publisher. This allows us to gain more immediate insight from the recommendations produced by the models.

2.2. Model

We build four models using a deep autoencoder for two models and a variational autoencoder for the other two models using the Keras framework. All the models we build have the same input of the user preference probability matrix. However, two models with differing autoencoders are created that utilize three latent variable matrices with each user's game names, game genres and reviews respectively. Every users information is separated by row in each matrix, with the ordering of rows being consistent across all matrices. The models output is a scoring matrix with the same dimension as the preference matrix S .




Figure 1. Deep Autoencoder model with Latent Matrices

The architecture of the deep autoencoder implemented can be seen in Figure 1. The user preference matrix passes through the encoder neural network consisting of two SELU dense layers that reduces the spatial size of the representation to 256. The other input matrices are each embedded, flattened, and then pass through a SELU dense layer that reduces the spatial size to 256. The two 256-size layers make

up the latent space of the autoencoder. The data in the latent space then passes through a dropout layer to avoid overfitting, and then passes through the decoder neural network that converts the hidden representation to one intermediate dense layer of 512 units. Finally, the data is converted back to its original dimension through a linear dense layer.




Figure 2. Variational Autoencoder model with Latent Matrices

The architecture of the variational autoencoder can be seen in Figure 2. In the variational autoencoder, we assume that the probability distribution $P(z|X)$, which is the probability of latent variable z given input X is Gaussian to learn its parameters. Instead of outputting a hidden representation z , the encoder outputs mean $\mu(X)$ and $\log \sigma(X)$. Next, we use the Gaussian reparameterization trick such that the operation

$$z = \mu(X) + \exp\left(\frac{\log \sigma(X)}{2}\right) \cdot \epsilon$$

where $\epsilon \sim N(0, 1)$, is equivalent to a Gaussian sampling because the sampling process does not have a gradient. Finally, the decoder $P(X|z)$ is similar to the decoder in the deep autoencoder with the exception that the activation function of the output layer is the non-linear sigmoid function.

2.3. Recommender

Once we have the reconstructed matrix as an output from the autoencoder, we set the entries of played games to 0 in order to not recommend games that the user had previously played. Finally, a recommender is built to present the top recommended games for a queried user based on the highest scoring games from the models output matrix.

3. Results

3.1. Loss and Mean Squared Error

The mean square loss function is used to train the deep autoencoder and the general loss function for the variational autoencoder is shown as

$$E[(X|z)] - KL[Q(z|X)|P(z)]$$

The first term for our equation, which is the construction loss, can be calculated by the binary cross entropy C, given

Table 1. Mean Squared Error after training

MODELS	LEARNING RATE	EPOCHS	MSE
DAE	0.00005	50	0.0018
DAE WITH LATENTS	0.00005	50	0.0026
VAE	0.0001	50	0.0056
VAE WITH LATENTS	0.00005	50	0.0082

input X and predicted probability $P(X|z)$:

$$C = \frac{1}{n} \sum_{i=1}^n [X_i \ln P(X_i|z) + (1 - X_i) \ln(1 - P(X_i|z))]$$

This is possible because the input for our model contains probability values between zero and one. The KL divergence term is implemented as:

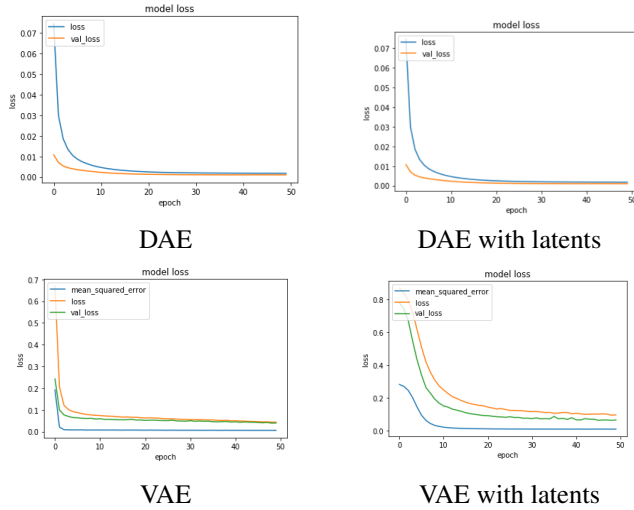
$$KL = 0.5 * \sum [\exp(\log \sigma) + \mu^2 - 1 - \log \sigma]$$

Although the two types of autoencoders use different loss functions to train, the Keras fit method can calculate the mean squared error (MSE) of the variational autoencoder while training. This metric is provided for equal comparison between the two autoencoders.

We experiment with the deep autoencoders (DAE) and the variational autoencoders (VAE) and produce four models. The MSE obtained after the final epoch of training from each experiment is shown in Table 1. In addition, plots of the loss curve from each experiment is shown in Table 2 where we see all the models converging at the end of their training. Because the DAE models were trained with the mean squared function, the loss shown in these graphs is the mean squared error of that model.

We see from Table 1 that the DAE models both with and without latent variables converge to a lower MSE than either VAE models even if those models initially had a lower MSE at the start of training. We also see that the MSE is higher for models using the latent matrices than without them. In Table 2, we see that most of the models start with high MSE then very quickly converge to low MSE. However, the VAE with latent variables shows a slower convergence, and at the end of training the validation loss does not plateau like the other models. This indicates that the latent matrices have more impact on the VAE models than the DAE models. With larger latent matrices, we can expect to see larger changes with how VAE models train than the DAE models.

Table 2. Plots of history loss during training



3.2. Recommendations

We review the information of user A from the dataset and the recommendations produced for user A to analyze the models' output. The history of played games from user A can be seen in Figure 3, while Figures 4-7 contain recommendations by the various models. The highest recommended game by each of the four models is the game "Garry's Mod". Additionally, the game "Saint Rows 2" appears in all four recommendations, while other games like "Half Life" and "Star Trek Online" appear in most of the lists. This indicates our models have similarities in feature extraction. We therefore look at the confidence score indicated by the output of the matrix. We see that DAE has the lowest score, indicating a weak recommendation, while the VAEs have a drastically higher score.

User played:	name	genres	publisher	developer	price
game_id	title				
790	22532	Counter-Strike: Global Offensive	Valve	Valve	14.99
8930	10345	Sid Meier's Civilization® 4	Firaxis Games, Aspyr (Mac, Linux)	2K Games, Aspyr (Mac, Linux)	29.99
1250	10000	Killing Floor	Tripwire Interactive	Tripwire Interactive	19.99
300	4733	Day of Defeat: Source	Valve	Valve	9.99
3590	4413	Plants vs. Zombies GOTY Edition	PopCap Games, Inc.	PopCap Games, Inc.	4.99
8190	3083	Just Cause 2	Avalanche Studios	Square Enix	14.99
8960	3561	Bushyland	Gambler Software	2K Games	19.99
6910	2685	Deus Ex: Game of the Year Edition	Ion Storm	Square Enix	9.99
8870	2084	Blackhawk Infinite	Irregular Games, Aspyr (Mac), Virtual Programm...	2K Games, Aspyr (Mac)	29.99
240	1853	Counter-Strike: Source	Valve	Valve	19.99

Figure 3. User A game history

score	name	genres	publisher	developer	price	
game_id						
4000	0.109021	Garry's Mod	['India', 'Simulation']	Farmponch Studios	Valve	9.99
4500	0.109003	S.T.A.L.K.E.R.: Shadow of Chernobyl	['Action', 'RPG']	GSC Game World	GSC Game World	19.99
9480	0.081378	Saints Row 2	['Action']	Valiant	Deep Silver	9.99
70	0.075740	Half-Life	['Action']	Valve	Valve	9.99
320	0.070106	Half-Life 2: Deathmatch	['Action']	Valve	Valve	4.99
6960	0.069595	Ultima: Blood Money	['Action']	Io-Interactive A/S	Io-Interactive A/S	5.99
8900	0.066595	EVE Online	['Action', 'Free to Play', 'Massively Multipla...']	COP	COP	0.00
2100	0.025235	Dark Messiah of Night Hap: Reac...	['Action', 'RPG']	Akshay Studios	Reinhardt	9.99
2600	0.046376	Vampire: The Masquerade - Bloodlines	['Action']	Travis Games	Activision	19.99
1520	0.044616	Counter-Strike: Source	['India', 'Strategy']	Introverson Software	Introverson Software	9.99

Figure 4. User A Recommendations by Deep Autoencoder

game_id	score	name	genres	publisher	developer	price
4000	0.395532	Garry's Mod	['India', 'Simulation']	Farmponch Studios	Valve	9.99
9480	0.143398	Saints Row 2	['Action']	Valiant	Deep Silver	9.99
4020	0.030239	Natural Selection 2	['Action', 'India', 'Strategy']	Unknown Worlds Entertainment	Unknown Worlds Entertainment	9.99
9450	0.083104	Warhammer® 40,000: Dawn of War® - Soulstorm	['Strategy']	Bolix Entertainment	BBIX	12.99
4760	0.080281	Rome: Total War® - Collection	['Strategy']	The Creative Assembly	BBIX	9.99
80	0.067138	Counter-Strike: Condition Zero	['Action']	Valve	Valve	9.99
4580	0.050905	Warhammer® 40,000: Dawn of War® - Dark Crusade	['Strategy']	Bolix Entertainment	BBIX	12.99
2100	0.050525	Dark Messiah of Night Hap: Reac...	['Action', 'RPG']	Akshay Studios	Reinhardt	9.99
4570	0.050038	Warhammer® 40,000: Dawn of War® - Game of the ...	['Action', 'RPG']	Bolix Entertainment	BBIX	12.99
8850	0.049570	Civilization IV: Beyond the Sword	['Strategy']	Paradox Games	2K Games	9.99

Figure 5. User A Recommendation by Deep Autoencoder with Latent Matrices

game_id	score	name	genres	publisher	developer	price
4000	0.561401	Garry's Mod	['India', 'Simulation']	Farmponch Studios	Valve	9.99
4000	0.108302	Star Wars: Battlefront 2 (Classic, 2005)	['Action']	Lucasfilm, LucasArts, Disney Interactive	Lucasfilm, LucasArts, Disney Interactive	9.99
4020	0.083901	Natural Selection 2	['Action', 'India', 'Strategy']	Unknown Worlds Entertainment	Unknown Worlds Entertainment	9.99
9450	0.081449	Rome: Total War® - Collection	['Free to Play', 'Massively Multiplayer', 'RPG']	The Creative Assembly	Perfect World Entertainment	0.00
4760	0.046987	Changeling Online	['Free to Play', 'Massively Multiplayer', 'RPG']	Cryptic Studios	Perfect World Entertainment	0.00
9480	0.043171	Warhammer® 40,000: Dawn of War® - Soulstorm	['Strategy']	Bolix Entertainment	BBIX	12.99
9420	0.043425	Supreme Commander: Forged Alliance	['Strategy']	Our Forged Games	TBS Media	14.99
70	0.042420	Half-Life	['Action']	Valve	Valve	9.99
9480	0.041423	Saints Row 2	['Action']	Valiant	Deep Silver	9.99

Figure 6. User A Recommendation by Variational Autoencoder

game_id	score	name	genres	publisher	developer	price
4000	0.529602	Garry's Mod	['India', 'Simulation']	Farmponch Studios	Valve	9.99
4000	0.096627	Star Wars: Battlefront 2 (Classic, 2005)	['Action']	Lucasfilm, LucasArts, Disney Interactive	Lucasfilm, LucasArts, Disney Interactive	9.99
30	0.050050	Half-Life	['Action']	Valve	Valve	9.99
4000	0.047799	Natural Selection 2	['Action', 'India', 'Strategy']	Unknown Worlds Entertainment	Unknown Worlds Entertainment	9.99
9600	0.041444	Star Trek Online	['Free to Play', 'Massively Multiplayer', 'RPG']	Cryptic Studios	Perfect World Entertainment	0.00
9480	0.041444	Saints Row 2	['Action']	Valiant	Deep Silver	9.99
320	0.031994	Half-Life 2: Deathmatch	['Action']	Valve	Valve	4.99
4760	0.031090	Rome: Total War® - Collection	['Strategy']	The Creative Assembly	BBIX	9.99
9420	0.031383	Supreme Commander: Forged Alliance	['Strategy']	Our Forged Games	TBS Media	14.99
9450	0.030629	Warhammer® 40,000: Dawn of War® - Soulstorm	['Strategy']	Bolix Entertainment	BBIX	12.99

Figure 7. User A Recommendation by Variational Autoencoder with Latent Matrices

4. Future Work

While we have implemented various high-performance autoencoders based on contemporary research, there are other types of implementations that we can consider in the future. Currently our VAE models assume a standard Gaussian prior. Another implementation consists of a high performing VAE model that assumes user-dependent priors defined as functions of user review text instead (Karamanolakis et al., 2018), which had modeled the trends within their dataset with higher accuracy. In addition, while observations were made on the experiments, more developed performance metrics could be implemented for a more direct comparison between the two types of autoencoders.

5. Conclusion

We create models of deep and variational autoencoders on a Steam dataset to recommend games to users. Both autoencoders implemented have success in recommendations that fit the types of games users had already played in the past. While the mean squared error of deep autoencoder models is lower than that of variational autoencoder models, we see a greater effect with latent matrices on the latter models than the former. Additionally, we note that the variational autoencoder models output notably higher confidence scores for recommendations, indicating that a user has a higher probability of enjoying the recommended game. With a more diverse and expansive dataset and greater

computational power for running the models, we expect to see an improvement in accuracy and confidence with the recommendations.

References

- Apurva Pathak, Kshitiz Gupta, J. M. Generating and personalizing bundle recommendations on steam, 2017.
- Doersch, C. Tutorial on variational autoencoders, 2016.
- Kang, W. C. and McAuley, J. Self-attentive sequential recommendation, 2018.
- Karamanolakis, G., Cherian, K. R., Narayan, A. R., Yuan, J., Tang, D., and Jebara, T. Item recommendation with variational autoencoders and heterogenous priors. 2018. doi: 10.1145/3270323.327032.
- Kingma, D. P. and Welling, M. Auto-encoding variational bayes, 2013.
- Kuchaiev, O. and Ginsburg, B. Training deep autoencoders for collaborative filtering, 2017.
- Liang, D., Krishnan, R. G., Hoffman, M. D., and Jebara, T. Variational autoencoders for collaborative filtering, 2018.
- Moussawi, A. Towards large scale training of autoencoders for collaborative filtering, 2018.
- Wan, M. and McAuley, J. Item recommendation on monotonic behavior chains, 2018.